

```
In [1]: # import general libraries for data manipulation
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from sklearn.metrics.pairwise import cosine_similarity
from scipy.sparse import csr_matrix
warnings.filterwarnings("ignore")
from surprise import Reader, Dataset, SVD, KNNBaseline
from surprise.model_selection import train_test_split
from surprise import accuracy
from collections import defaultdict
```

```
c:\Users\rauna\Anaconda3\envs\scaler\lib\site-packages\numpy\_distributor_init.py:
30: UserWarning: loaded more than 1 DLL from .libs:
c:\Users\rauna\Anaconda3\envs\scaler\lib\site-packages\numpy\.libs\libopenblas.FB5
AE2TYXYH2IJRDKGDGQ3XBLKTF43H.gfortran-win_amd64.dll
c:\Users\rauna\Anaconda3\envs\scaler\lib\site-packages\numpy\.libs\libopenblas.GK7
GX5KEQ4F6UY03P26ULGBQYHGQ07J4.gfortran-win_amd64.dll
warnings.warn("loaded more than 1 DLL from .libs:")
```

## Loading Data

```
In [2]: movies = pd.read_fwf('ZEE-data/zee-movies.dat', encoding='ISO-8859-1')
ratings = pd.read_fwf('ZEE-data/zee-ratings.dat', encoding='ISO-8859-1')
users = pd.read_fwf('ZEE-data/zee-users.dat', encoding='ISO-8859-1')
```

## Preprocessing

```
In [3]: movies.drop(columns=['Unnamed: 1', 'Unnamed: 2'], inplace=True)
movies = movies['Movie ID::Title::Genres'].str.split(':', expand=True)
movies.columns = ['Movie ID', 'Title', 'Genres']
display(movies.head())
print(movies.shape)
```

	Movie ID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

(3883, 3)

```
In [4]: # do same for ratings
ratings = ratings['UserID::MovieID::Rating::Timestamp'].str.split(':', expand=True)
```

```
ratings.columns = ['User ID', 'Movie ID', 'Rating', 'Timestamp']
ratings.drop(columns=['Timestamp'], inplace=True)
display(ratings.head())
print(ratings.shape)
```

	User ID	Movie ID	Rating
0	1	1193	5
1	1	661	3
2	1	914	3
3	1	3408	4
4	1	2355	5

(1000209, 3)

```
In [5]: users = users['UserID::Gender::Age::Occupation::Zip-code'].str.split(':', expand=True)
users.columns = ['User ID', 'Gender', 'Age', 'Occupation', 'Zip-code']
users.drop(columns=['Zip-code'], inplace=True)
display(users.head())
print(users.shape)
```

	User ID	Gender	Age	Occupation
0	1	F	1	10
1	2	M	56	16
2	3	M	25	15
3	4	M	45	7
4	5	M	25	20

(6040, 4)

```
In [6]: # merge the dataframes
df = pd.merge(movies, ratings, on='Movie ID')
df = pd.merge(df, users, on='User ID')
display(df.head())
print(df.shape)
```

	Movie ID	Title	Genres	User ID	Rating	Gender	Age	Occupation
0	1	Toy Story (1995)	Animation Children's Comedy	1	5	F	1	1
1	48	Pocahontas (1995)	Animation Children's Musical Romance	1	5	F	1	1
2	150	Apollo 13 (1995)	Drama	1	5	F	1	1
3	260	Star Wars: Episode IV - A New Hope (1977)	Action Adventure Fantasy	1	4	F	1	1
4	527	Schindler's List (1993)	Drama War	1	5	F	1	1

(1000209, 8)

## EDA

```
In [7]: df.shape, users.shape, movies.shape
```

```
Out[7]: ((1000209, 8), (6040, 4), (3883, 3))
```

- We have 1 million ratings from 6000 users on 4000 movies

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1000209 entries, 0 to 1000208
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Movie ID        1000209 non-null object
1   Title           1000209 non-null object
2   Genres          996144 non-null object
3   User ID         1000209 non-null object
4   Rating          1000209 non-null object
5   Gender          1000209 non-null object
6   Age            1000209 non-null object
7   Occupation      1000209 non-null object
dtypes: object(8)
memory usage: 68.7+ MB
```

- Genres has some missing values

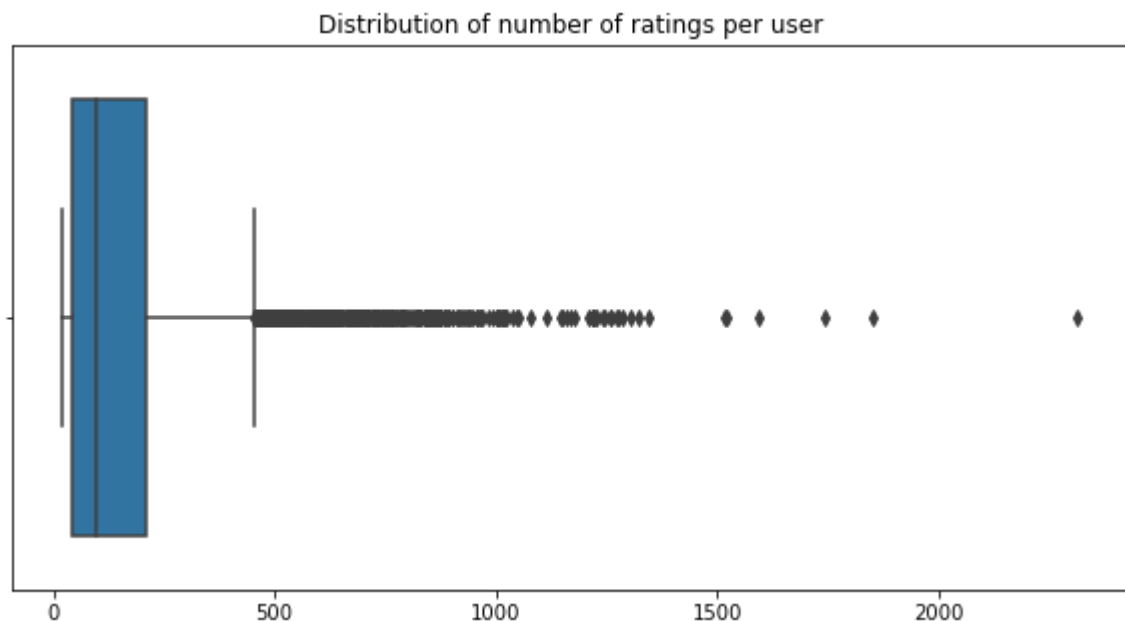
```
In [9]: for c in ['Rating', 'Age', 'Movie ID', 'User ID']:
        df[c] = df[c].astype('int')
```

```
In [10]: # data checks
display(df['User ID'].agg(['min', 'max']))
display(df['Movie ID'].agg(['min', 'max']))
display(df['Rating'].agg(['min', 'max']))
print(f"Minimum ratings per user : {df.groupby('User ID').size().min()}")

min      1
max    6040
Name: User ID, dtype: int32
min      1
max    3952
Name: Movie ID, dtype: int32
min      1
max      5
Name: Rating, dtype: int32
Minimum ratings per user : 20
```

- All the above checks are in line with data description

```
In [11]: # Outlier detection
plt.figure(figsize=(10,5))
sns.boxplot(df.groupby('User ID').size())
plt.title('Distribution of number of ratings per user')
plt.show()
```



```
In [12]: # write function to remove outliers using IQR
def remove_outliers(df, col, k=1.5):
    q1 = df[col].quantile(0.25)
    q3 = df[col].quantile(0.75)
    iqr = q3 - q1
    lower_bound = q1 - (k * iqr)
    upper_bound = q3 + (k * iqr)
    return df[(df[col] > lower_bound) & (df[col] < upper_bound)]
```

```
In [13]: valid_users = remove_outliers(pd.DataFrame(df.groupby('User ID').size()), col=0, k=
```

```
In [14]: print(df['User ID'].unique(), valid_users.unique(), df['Movie ID'].unique())
df = df[df['User ID'].isin(valid_users)]
print(df['User ID'].unique(), valid_users.unique(), df['Movie ID'].unique())

6040 5564 3706
5564 5564 3613
```

## Feature Creation

```
In [15]: # create a new column for year
df['ReleaseYear'] = df['Title'].str.extract('.*\((.*)\).*', expand=True)

# create new feature for total number of words in title
df['TitleWordCount'] = df['Title'].str.split().str.len()

# create new feature for number of genres
df['GenreCount'] = df['Genres'].str.count('\|') + 1
df['GenreCount'].fillna(0, inplace=True)
```

```
In [16]: df.head()
```

```
Out[16]:
```

	Movie ID	Title	Genres	User ID	Rating	Gender	Age	Occupation
0	1	Toy Story (1995)	Animation Children's Comedy	1	5	F	1	1
1	48	Pocahontas (1995)	Animation Children's Musical Romance	1	5	F	1	1
2	150	Apollo 13 (1995)	Drama	1	5	F	1	1
3	260	Star Wars: Episode IV - A New Hope (1977)	Action Adventure Fantasy	1	4	F	1	1
4	527	Schindler's List (1993)	Drama War	1	5	F	1	1

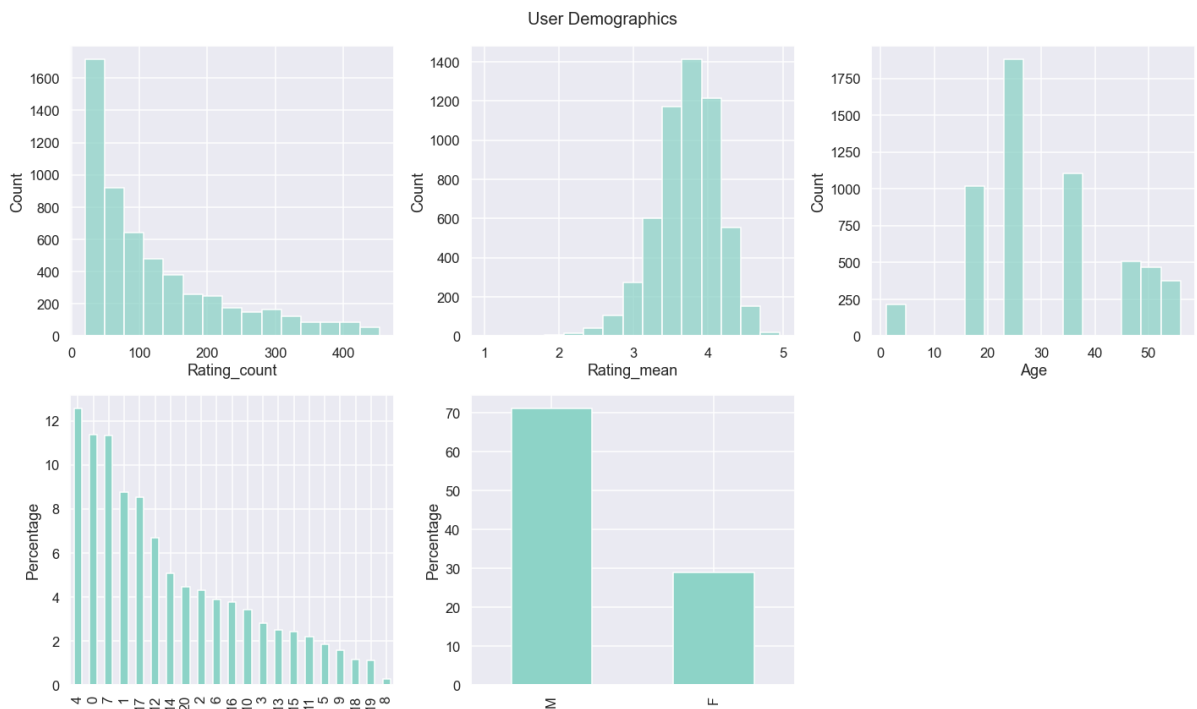
```
In [17]: df_grouped = df.groupby('User ID').agg({'Rating': ['count', 'mean'], 'Gender': 'first'})
df_grouped.columns = ['_'.join(col).strip() for col in df_grouped.columns.values]
df_grouped.reset_index(inplace=True)
df_grouped.rename(columns={'Gender_first': 'Gender', 'Age_first': 'Age', 'Occupation_first': 'Occupation'})
display(df_grouped.head())
print(df_grouped.shape)
```

	User ID	Rating_count	Rating_mean	Gender	Age	Occupation
0	1	53	4.188679	F	1	10
1	2	129	3.713178	M	56	16
2	3	51	3.901961	M	25	15
3	4	21	4.190476	M	45	7
4	5	198	3.146465	M	25	20

(5564, 6)

```
In [18]: sns.set_style('darkgrid')
sns.set_context('talk')
sns.set_palette('Set3')
plt.rcParams['figure.autolayout'] = True


fig, axes = plt.subplots(2, 3, figsize=(20, 12))
# set title
fig.suptitle('User Demographics', fontsize=20)
for e, c in enumerate(['Rating_count', 'Rating_mean', 'Age', 'Occupation', 'Gender']):
    if df_grouped[c].dtype == 'object':
        # plot countplot as percentage
        (df_grouped[c].value_counts(normalize=True)*100).plot(kind='bar', ax=axes[e//3, e%3])
        axes[e//3, e%3].set_ylabel('Percentage')
    else:
        sns.histplot(df_grouped[c], ax=axes[e//3, e%3], bins=15)
# remove last subplot
fig.delaxes(axes[1, 2])
plt.show()
```

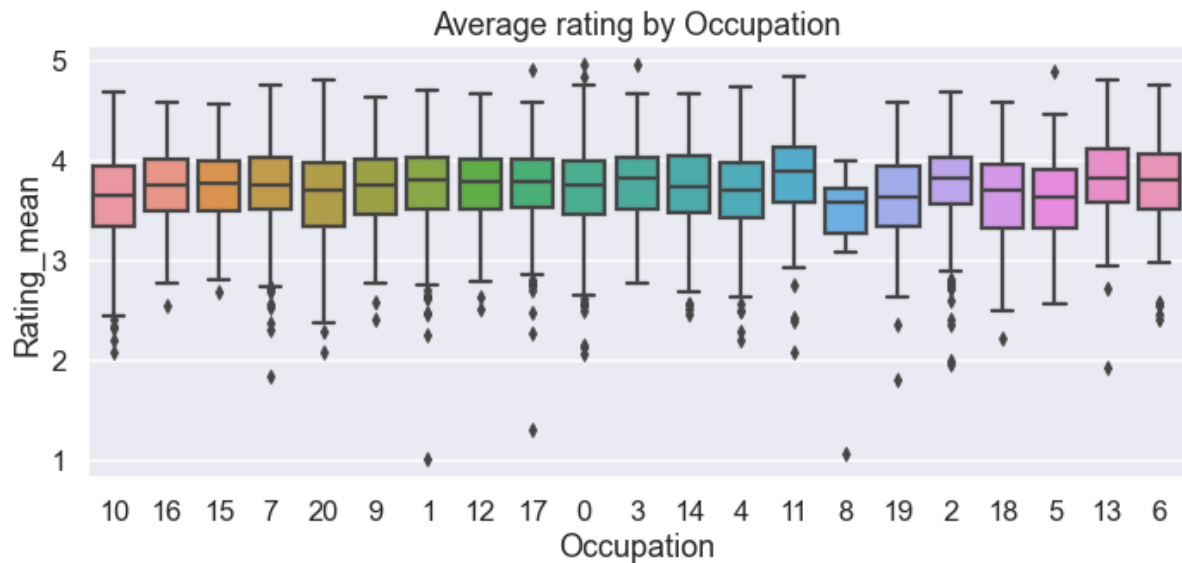


- Users have an exponential distribution when it comes to # of movies rated
- Users usually rate movies between 3 and 4

- Primary user base is in 20s
- College/Grad Student makes up the highest of our user base ~ 12%
- 70 % of users are male

```
In [19]: for e,c in enumerate(['Age', 'Gender', 'Occupation']):
plt.figure(figsize=(10,5))
sns.boxplot(x=c, y='Rating_mean', data=df_grouped)
plt.title(f'Average rating by {c}')
plt.show()
```





- There is a slight positive relationship between age and average user rating showing older people are less critical
- Gender doesn't have much difference in ratings
- Occupation 8 ie Farmers rate the lowest on average

```
In [20]: # Creating a pivot table of movie titles & user id and imputing the NaN values with
user_movie_ratings = df.pivot_table(index='User ID', columns='Movie ID', values='Ra
display(user_movie_ratings.head())
print(user_movie_ratings.shape)
```

Movie ID	1	2	3	4	5	6	7	8	9	10	...	3943	3944	3945	3946	3947	3948	:
User ID																		
1	5.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
5	0.0	0.0	0.0	0.0	0.0	2.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	

5 rows × 3613 columns

(5564, 3613)

```
In [21]: print(f"Empty percentage : {round((user_movie_ratings == 0).mean().mean() * 100,1)}")
Empty percentage : 96.7 %
```

## Recommendation System



```
In [22]: df.head()
```

```
Out[22]:
```

	Movie ID	Title	Genres	User ID	Rating	Gender	Age	Occupation
0	1	Toy Story (1995)	Animation Children's Comedy	1	5	F	1	1
1	48	Pocahontas (1995)	Animation Children's Musical Romance	1	5	F	1	1
2	150	Apollo 13 (1995)	Drama	1	5	F	1	1
3	260	Star Wars: Episode IV - A New Hope (1977)	Action Adventure Fantasy	1	4	F	1	1
4	527	Schindler's List (1993)	Drama War	1	5	F	1	1

```
In [23]: # Setting up
reader = Reader(rating_scale=(1, 5))
data = Dataset.load_from_df(df[['User ID', 'Movie ID', 'Rating']], reader)
trainset, testset = train_test_split(data, test_size=0.2, random_state=42)
```

```
In [24]: def build_rec_sys(algo):
        algo.fit(trainset)
        predictions = algo.test(testset)
        print(f"Test RMSE : {round(accuracy.rmse(predictions),2)} ,Test MAE: {round(accuracy.mae(predictions),2)}")
        return algo, predictions
```

## Using Pearson Correlation

```
In [25]: model_1 , predictions_1 = build_rec_sys(KNNBaseline(sim_options={'name': 'pearson_baseline'}))

Estimating biases using als...
Computing the pearson_baseline similarity matrix...
Done computing similarity matrix.
RMSE: 0.8820
MAE: 0.6894
Test RMSE : 0.88 ,Test MAE: 0.69
```

```
In [26]: movie_dict = dict(zip(df['Movie ID'], df['Title']))

def get_top_n(predictions, movie_dict, n=10):
    top_n = defaultdict(list)
    for uid, iid, true_r, est, _ in predictions:
        top_n[uid].append((movie_dict[iid], est ))
    for uid, user_ratings in top_n.items():
        user_ratings.sort(key=lambda x : x[1], reverse=True)
        top_n[uid] = user_ratings[:n]
    return top_n
```

```
In [27]: top_n_1 = get_top_n(predictions_1, movie_dict, n=10)
```

## Top 10 recommendations for user 1

```
In [28]: user_1 = df[df['User ID'] == 1]
user_1 = user_1[['User ID', 'Movie ID', 'Title', 'Rating']]
print('\nMovies liked by user 1: ')
display(user_1.sort_values('Rating', ascending=False).head(10))
print('\nTop 10 recommendations for user 1')
for e,i in enumerate(top_n_1[1]):
    print(f"{e+1}. {i[0]} with predicted rating {round(i[1],2)}")
```

Movies liked by user 1:

	User ID	Movie ID	Title	Rating
0	1	1	Toy Story (1995)	5
17	1	1022	Cinderella (1950)	5
49	1	3105	Awakenings (1990)	5
47	1	2804	Christmas Story, A (1983)	5
40	1	2355	Bug's Life, A (1998)	5
36	1	2028	Saving Private Ryan (1998)	5
33	1	1961	Rain Man (1988)	5
31	1	1836	Last Days of Disco, The (1998)	5
27	1	1287	Ben-Hur (1959)	5
1	1	48	Pocahontas (1995)	5

Top 10 recommendations for user 1

1. Wizard of Oz, The (1939) with predicted rating 4.72
2. Ben-Hur (1959) with predicted rating 4.6
3. Star Wars: Episode IV - A New Hope (1977) with predicted rating 4.55
4. E.T. the Extra-Terrestrial (1982) with predicted rating 4.45
5. Dead Poets Society (1989) with predicted rating 4.39
6. Beauty and the Beast (1991) with predicted rating 4.34
7. Tarzan (1999) with predicted rating 4.1
8. Secret Garden, The (1993) with predicted rating 3.91
9. Hunchback of Notre Dame, The (1996) with predicted rating 3.8
10. Antz (1998) with predicted rating 3.59

## Top 10 movies similar to Toy Story (1995)

```
In [38]: toy_story_raw_id = df[df['Title'] == 'Toy Story (1995)']['Movie ID'].iloc[0]
```

```
In [30]: [movie_dict[trainset.to_raw_iid(i)] for i in model_1.get_neighbors(trainset.to_inne
```

```
Out[30]: ['Toy Story 2 (1999)',
          "Bug's Life, A (1998)",
          'Aladdin (1992)',
          'Beauty and the Beast (1991)',
          'Honey, I Shrunk the Kids (1989)',
          'Muppet Movie, The (1979)',
          'Babe (1995)',
          'Lion King, The (1994)',
          'Tarzan (1999)',
          'Cinderella (1950)']
```

## Using Cosine Similarity

### Item Item Similarity

```
In [31]: model_2 , predictions_2 = build_rec_sys(KNNBaseline(sim_options={'name': 'cosine',
Estimating biases using als...
Computing the cosine similarity matrix...
Done computing similarity matrix.
RMSE: 0.9112
MAE: 0.7167
Test RMSE : 0.91 ,Test MAE: 0.72

In [36]: item_similarity_matrix = pd.DataFrame(model_2.compute_similarities())
item_similarity_matrix.index = item_similarity_matrix.index.map(lambda x : movie_di
item_similarity_matrix.columns = item_similarity_matrix.index
item_similarity_matrix.head()

Computing the cosine similarity matrix...
Done computing similarity matrix.
```

Out[36]:

	First Wives Club, The (1996)	Sphere (1998)	L.A. Confidential (1997)	Rocky Horror Picture Show, The (1975)	Lone Star (1996)	Stand by Me (1986)	On the Waterfront (1954)	Chinatown (1974)
First Wives Club, The (1996)	1.000000	0.926677	0.957309	0.923407	0.962500	0.949004	0.954030	0.94070
Sphere (1998)	0.926677	1.000000	0.927770	0.875052	0.926874	0.934048	0.879211	0.90605
L.A. Confidential (1997)	0.957309	0.927770	1.000000	0.929263	0.970224	0.969329	0.974220	0.98159
Rocky Horror Picture Show, The (1975)	0.923407	0.875052	0.929263	1.000000	0.936119	0.917318	0.910525	0.93178
Lone Star (1996)	0.962500	0.926874	0.970224	0.936119	1.000000	0.963328	0.980775	0.97718

5 rows × 3575 columns

## Top 10 movies similar to Toy Story (1995)

```
In [41]: [movie_dict[trainset.to_raw_iid(i)] for i in model_2.get_neighbors(trainset.to_inne
```

```
Out[41]: ['Rendezvous in Paris (Rendez-vous de Paris, Les) (1995)',  
'Buck and the Preacher (1972)',  
'Bad Moon (1996)',  
'Ape, The (1940)',  
'Solas (1999)',  
'Solar Crisis (1993)',  
'Story of G.I. Joe, The (1945)',  
'Little Lord Fauntleroy (1936)',  
'Don't Look in the Basement! (1973)',  
'Saltmen of Tibet, The (1997)']
```

## User User Similarity

```
In [42]: model_3 , predictions_3 = build_rec_sys(KNNBaseline(sim_options={'name': 'cosine',  
Estimating biases using als...  
Computing the cosine similarity matrix...  
Done computing similarity matrix.  
RMSE: 0.9165  
MAE: 0.7232  
Test RMSE : 0.92 ,Test MAE: 0.72
```

```
In [43]: user_similarity_matrix = pd.DataFrame(model_3.compute_similarities())
```

```
user_similarity_matrix.head()
```

```
Computing the cosine similarity matrix...  
Done computing similarity matrix.
```

```
Out[43]:
```

	0	1	2	3	4	5	6	7	8	
0	1.000000	0.991837	0.993809	0.992171	0.972677	0.957209	0.962583	0.938436	0.934335	0.9676
1	0.991837	1.000000	0.994850	0.967989	0.913393	0.993610	0.966842	0.936390	0.962720	0.8734
2	0.993809	0.994850	1.000000	0.994536	0.970429	0.990947	0.991408	0.972369	0.947534	0.9902
3	0.992171	0.967989	0.994536	1.000000	0.942957	0.987541	0.972629	0.968922	1.000000	0.7889
4	0.972677	0.913393	0.970429	0.942957	1.000000	0.968521	0.963240	0.944295	0.956269	0.9429

5 rows × 5564 columns

## Top 10 recommendations for user 1

```
In [45]: top_n_3 = get_top_n(predictions_3, movie_dict, n=10)
```

```
In [47]: user_1 = df[df['User ID'] == 1]  
user_1 = user_1[['User ID', 'Movie ID', 'Title', 'Rating']]  
print('\nMovies liked by user 1: ')  
display(user_1.sort_values('Rating', ascending=False).head(10))  
print('\nTop 10 recommendations for user 1')  
for e,i in enumerate(top_n_3[1]):  
    print(f"{e+1}. {i[0]} with predicted rating {round(i[1],2)}")
```

Movies liked by user 1:

	User ID	Movie ID	Title	Rating
0	1	1	Toy Story (1995)	5
17	1	1022	Cinderella (1950)	5
49	1	3105	Awakenings (1990)	5
47	1	2804	Christmas Story, A (1983)	5
40	1	2355	Bug's Life, A (1998)	5
36	1	2028	Saving Private Ryan (1998)	5
33	1	1961	Rain Man (1988)	5
31	1	1836	Last Days of Disco, The (1998)	5
27	1	1287	Ben-Hur (1959)	5
1	1	48	Pocahontas (1995)	5

Top 10 recommendations for user 1

1. Star Wars: Episode IV - A New Hope (1977) with predicted rating 4.53
2. Wizard of Oz, The (1939) with predicted rating 4.51
3. Ben-Hur (1959) with predicted rating 4.25
4. Dead Poets Society (1989) with predicted rating 4.24
5. Beauty and the Beast (1991) with predicted rating 4.17
6. Secret Garden, The (1993) with predicted rating 4.02
7. Tarzan (1999) with predicted rating 3.96
8. E.T. the Extra-Terrestrial (1982) with predicted rating 3.77
9. Antz (1998) with predicted rating 3.7
10. Hunchback of Notre Dame, The (1996) with predicted rating 3.64

## Using Matrix Decomposition

```
In [52]: model_4 , predictions_4 = build_rec_sys(SVD(n_factors=4, n_epochs=100))
```

RMSE: 0.8881

MAE: 0.6934

Test RMSE : 0.89 ,Test MAE: 0.69

## Top 10 recommendations for user 1

```
In [54]: top_n_4 = get_top_n(predictions_4, movie_dict, n=10)
```

```
In [55]: user_1 = df[df['User ID'] == 1]
user_1 = user_1[['User ID', 'Movie ID', 'Title', 'Rating']]
print('\nMovies liked by user 1: ')
display(user_1.sort_values('Rating', ascending=False).head(10))
print('\nTop 10 recommendations for user 1')
for e,i in enumerate(top_n_4[1]):
    print(f"{e+1}. {i[0]} with predicted rating {round(i[1],2)}")
```

Movies liked by user 1:

	User ID	Movie ID	Title	Rating
0	1	1	Toy Story (1995)	5
17	1	1022	Cinderella (1950)	5
49	1	3105	Awakenings (1990)	5
47	1	2804	Christmas Story, A (1983)	5
40	1	2355	Bug's Life, A (1998)	5
36	1	2028	Saving Private Ryan (1998)	5
33	1	1961	Rain Man (1988)	5
31	1	1836	Last Days of Disco, The (1998)	5
27	1	1287	Ben-Hur (1959)	5
1	1	48	Pocahontas (1995)	5

Top 10 recommendations for user 1

1. Wizard of Oz, The (1939) with predicted rating 4.61
2. Star Wars: Episode IV - A New Hope (1977) with predicted rating 4.5
3. Ben-Hur (1959) with predicted rating 4.41
4. E.T. the Extra-Terrestrial (1982) with predicted rating 4.37
5. Dead Poets Society (1989) with predicted rating 4.34
6. Beauty and the Beast (1991) with predicted rating 4.28
7. Tarzan (1999) with predicted rating 4.06
8. Secret Garden, The (1993) with predicted rating 3.99
9. Antz (1998) with predicted rating 3.74
10. Hunchback of Notre Dame, The (1996) with predicted rating 3.68

## Embeddings

```
In [63]: user_embeddings = pd.DataFrame(model_4.pu,
                                         index=[trainset.to_raw_uid(i) for i in range(trainset.n
                                         columns=['Embedding_1', 'Embedding_2', 'Embedding_3', '

items_embeddings = pd.DataFrame(model_4.qi,
                                index=[movie_dict[trainset.to_raw_iid(i)] for i in rang
                                columns=['Embedding_1', 'Embedding_2', 'Embedding_3', '
```

```
In [64]: print('User Embeddings : ')
display(user_embeddings.head(2))
print('\n\nItem Embeddings : ')
display(items_embeddings.head(2))
```

User Embeddings :

	Embedding_1	Embedding_2	Embedding_3	Embedding_4
<b>96</b>	-0.449350	-0.053914	-0.329087	-0.260830
<b>2101</b>	0.388652	-0.174531	0.033435	0.921491

Item Embeddings :

	Embedding_1	Embedding_2	Embedding_3	Embedding_4
<b>First Wives Club, The (1996)</b>	-0.501219	-0.191519	-0.594369	-0.474218
<b>Sphere (1998)</b>	-0.091231	0.331609	-0.044451	-0.739272

```
In [65]: user_similarity_cosine = pd.DataFrame(cosine_similarity(user_embeddings), index=use
item_similarity_cosine = pd.DataFrame(cosine_similarity(items_embeddings), index=it
```

```
In [66]: print('\nUser Similarity Matrix : ')
display(user_similarity_cosine.head(2))
print('\n\nItem Similarity Matrix : ')
display(item_similarity_cosine.head(2))
```

User Similarity Matrix :

	96	2101	5846	2284	1984	1699	2185	5699	1964
96	1.000000	-0.664299	0.310701	-0.620868	0.354540	0.081156	0.230757	-0.713970	0.523544
2101	-0.664299	1.000000	0.336351	0.841520	-0.393141	0.626136	0.523050	0.979375	-0.038011

2 rows × 5564 columns

Item Similarity Matrix :

	First Wives Club, The (1996)	Sphere (1998)	L.A. Confidential (1997)	Rocky Horror Picture Show, The (1975)	Lone Star (1996)	Stand by Me (1986)	On the Waterfront (1954)	Chinatown (1974)
First Wives Club, The (1996)	1.000000	0.472701	-0.587955	-0.438781	0.134283	-0.484963	-0.233838	-0.466187
Sphere (1998)	0.472701	1.000000	-0.943650	-0.429908	-0.653336	-0.676069	-0.947037	-0.787606

2 rows × 3575 columns

## Top 10 movies similar to Toy Story (1995)

```
In [71]: item_similarity_cosine['Toy Story (1995)'].sort_values(ascending=False).head(11)
```

```
Out[71]: Toy Story (1995)          1.000000
Jungle Book, The (1967)          0.996004
Splash (1984)                    0.991048
E.T. the Extra-Terrestrial (1982) 0.989616
Muppet Christmas Carol, The (1992) 0.984511
Sleeping Beauty (1959)           0.981426
Old Yeller (1957)                0.981204
Big (1988)                       0.974407
Oliver! (1968)                   0.968310
Ayn Rand: A Sense of Life (1997)  0.967853
Rescuers, The (1977)             0.966946
Name: Toy Story (1995), dtype: float64
```

## Embeddings Visulization with d=2

```
In [72]: model_5 , predictions_5 = build_rec_sys(SVD(n_factors=2, n_epochs=200))
```

```
RMSE: 0.8889
MAE: 0.6964
Test RMSE : 0.89 ,Test MAE: 0.7
```

```
In [159... user_embeddings = pd.DataFrame(model_5.pu,
                                index=[trainset.to_raw_uid(i) for i in range(trainset.n
```



```

        columns=['Embedding_1', 'Embedding_2'])

items_embeddings = pd.DataFrame(model_5.qi,
                                index=[movie_dict[trainset.to_raw_iid(i)] for i in range(
                                    columns=['Embedding_1', 'Embedding_2'])

```

```

In [160... # merge item embeddings with movie data
items_embeddings = movies.merge(items_embeddings, left_on='Title', right_index=True)

# extract year from title
items_embeddings['year'] = items_embeddings['Title'].str.extract('.*\((.*)\).*', ex
items_embeddings['year'] = pd.to_numeric(items_embeddings['year'], errors='coerce')
items_embeddings['released_after_1995'] = (items_embeddings['year'] > 1995)

# extract number of genres from Genre
items_embeddings['genre_count'] = items_embeddings['Genres'].str.count('\|') + 1
items_embeddings['multiple_genre'] = (items_embeddings['genre_count'] > 1)

display(items_embeddings.head(2))

```

	Movie ID	Title	Genres	Embedding_1	Embedding_2	year	released_after
0	1	Toy Story (1995)	Animation Children's Comedy	0.385643	0.232034	1995.0	
1	2	Jumanji (1995)	Adventure Children's Fantasy	0.605638	-0.419186	1995.0	

```

In [165... # merge user embeddings with user data
users['User ID'] = users['User ID'].astype('int')
users['Age'] = users['Age'].astype('int')
user_embeddings = users.merge(user_embeddings, left_on='User ID', right_index=True)
user_embeddings['older_than_30'] = (user_embeddings['Age'] > 30)

display(user_embeddings.head(2))

```

	User ID	Gender	Age	Occupation	Embedding_1	Embedding_2	older_than_30
0	1	F	1	10	0.573303	-0.115888	False
1	2	M	56	16	0.854225	-0.092019	True

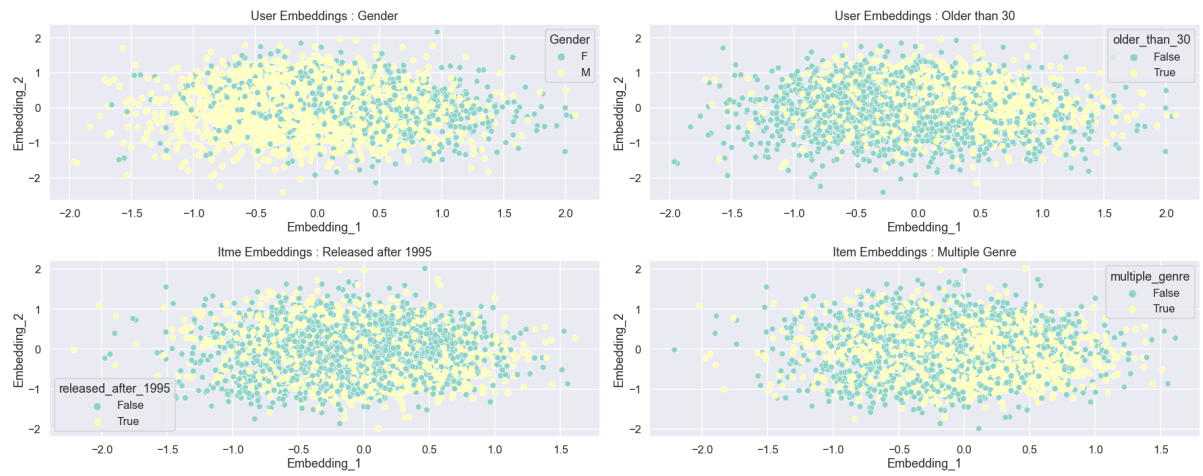
```

In [166... fig,ax = plt.subplots(2,2,figsize=(25, 10))
sns.scatterplot(x='Embedding_1', y='Embedding_2', data=user_embeddings, hue='Gender')
ax[0,0].set_title('User Embeddings : Gender')
sns.scatterplot(x='Embedding_1', y='Embedding_2', data=user_embeddings, hue='older_than_30')
ax[0,1].set_title('User Embeddings : Older than 30')

sns.scatterplot(x='Embedding_1', y='Embedding_2', data=items_embeddings, hue='released_after_1995')
ax[1,0].set_title('Item Embeddings : Released after 1995')
sns.scatterplot(x='Embedding_1', y='Embedding_2', data=items_embeddings, hue='multiple_genre')
ax[1,1].set_title('Item Embeddings : Multiple Genre')

```

```
Out[166]: Text(0.5, 1.0, 'Item Embeddings : Multiple Genre')
```



- No correlation between the embeddings
- Embeddings have captured gender differences (female top right, male bottom left)